

# **High Performance Computing, Cloud Computing**

**JSC 370: Data Science II**


# What is HPC

High Performance Computing (HPC) can relate to any of the following:

- **Parallel computing**, i.e. using multiple resources (could be threads, cores, nodes, etc.) simultaneously to complete a task.
- **Big data** working with large datasets (in/out-of-memory).


We will mostly focus on parallel computing.

# Serial computation



Here we are using a single core. The function is applied one element at a time, leaving the other 3 cores without usage.

# Parallel computation



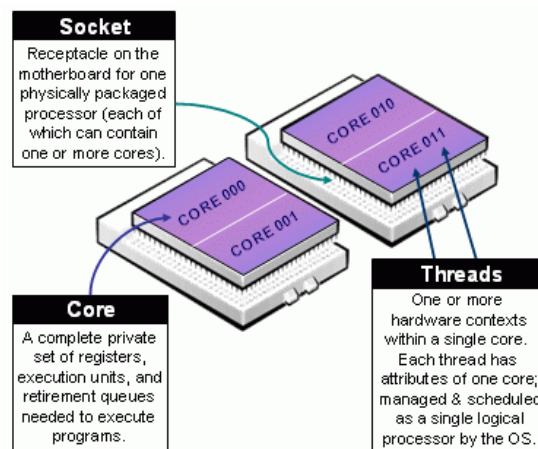
In this more intelligent way of computation, we are taking full advantage of our computer by using all 4 cores at the same time. This will translate in a reduced computation time which, in the case of complicated/long calculations, can be an important speed gain.

# Parallel computing: Hardware

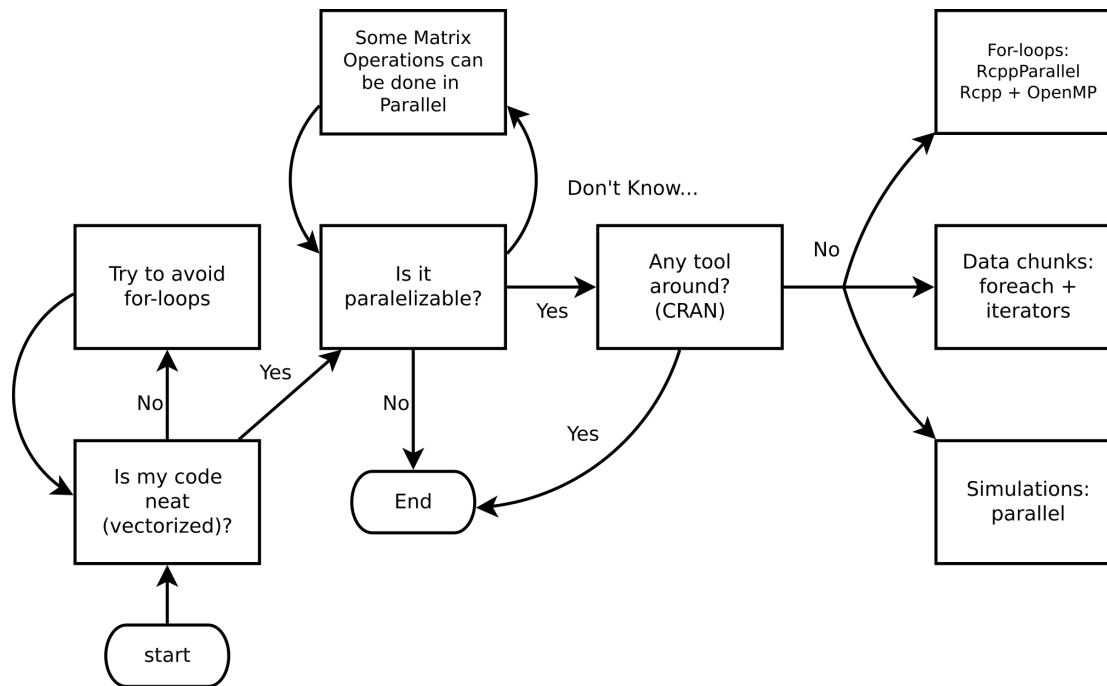
When it comes to parallel computing, there are several ways (levels) in which we can speed up our analysis. From the bottom up:

- **Thread level SIMD instructions**: In most modern processors support some level of what is called vectorization, this is, applying a single (same) instruction to streams of data, for example: adding vector A and B.
- **Hyper-Threading Technology** (HTT): Intel's hyper-threading generates a virtual partition of a single core (processor) which, while not equivalent to having multiple physical threads, does speedup things.
- **Multi-core processor**: Most modern CPUs (Central Processing Unit) have two or more physical cores. A typical laptop computer has about 8 cores.
- **General-Purpose Computing on Graphics Processing Unit** (GP-GPU): While modern CPUs have a couple of dozens of cores, GPUs can hold thousands of those. Designed for image processing, there's an increasing use of GPUs as an alternative of CPUs for scientific computing.
- **High-Performance Computing Cluster** (HPC): A collection of computing nodes that are interconnected using a fast Ethernet network.
- **Grid Computing**: A collection of loosely interconnected machines that may or may not be in the same physical place, for example: HTCondor clusters.

# Parallel computing: CPU components



Taxonomy of CPUs (Downloaded from [https://slurm.schedmd.com/mc\\_support.html](https://slurm.schedmd.com/mc_support.html))



Ask yourself these questions before jumping into HPC!

# Parallel computing in R

While there are several ways to do parallel computing in R (just take a look at the [High-Performance Computing Task View](#)), we'll focus on the following R-packages for **explicit parallelism**

Some examples:

- [parallel](#): R package that provides '[s]upport for parallel computation, including random-number generation'.
- [foreach](#): R package for 'general iteration over elements' in parallel fashion.
- [future](#): '[A] lightweight and unified Future API for sequential and parallel processing of R expression via futures.'  
(won't cover here)

Implicit parallelism, on the other hand, are out-of-the-box tools that allow the programmer not to worry about parallelization, e.g. such as [gpuR](#) for Matrix manipulation using GPU, [tensorflow](#)

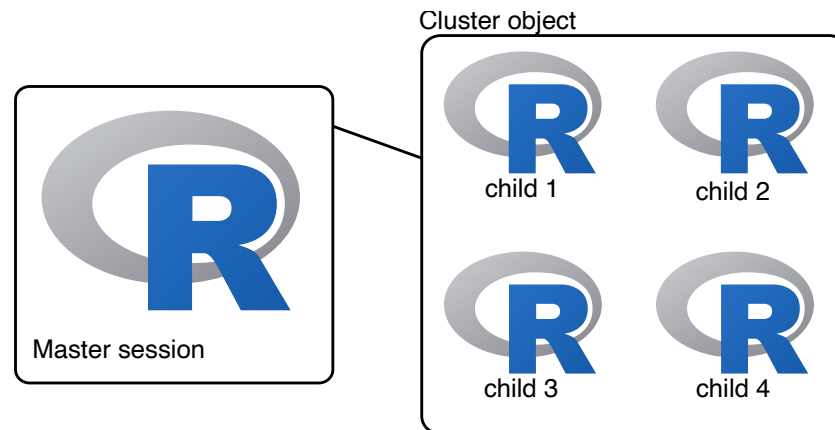


And there's also a more advanced set of options

- [Rcpp](#) + [OpenMP](#): [Rcpp](#) is an R package for integrating R with C++, and OpenMP is a library for high-level parallelism for C/C++ and Fortran.
- A ton of other type of resources, notably the tools for working with batch schedulers such as Slurm, HTCondor, etc.

# The parallel package

- Based on the `snow` and `multicore` R Packages.
- Explicit parallelism.
- Simple yet powerful idea: Parallel computing as multiple R sessions.
- Clusters can be made of both local and remote sessions
- Multiple types of cluster: `PSOCK`, `Fork`, `MPI`, etc.



# Parallel workflow

(Usually) We do the following:

1. Create a PSOCK/FORK (or other) cluster using `makePSOCKCluster`/`makeForkCluster` (or `makeCluster`)
2. Copy/prepare each R session (if you are using a PSOCK cluster):
  - a. Copy objects with `clusterExport`
  - b. Pass expressions with `clusterEvalQ`
  - c. Set a seed
3. Do your call: `parApply`, `parLapply`, etc.
4. Stop the cluster with `clusterStop`

# Ex 1: Hello world!

```
# 1. CREATING A CLUSTER
library(parallel)
cl <- makePSOCKcluster(4)
x <- 20

# 2. PREPARING THE CLUSTER
clusterSetRNGStream(cl, 123) # Equivalent to `set.seed(123)`
clusterExport(cl, "x")

# 3. DO YOUR CALL
clusterEvalQ(cl, {
  paste0("Hello from process #", Sys.getpid(), ". I see x and it is equal to ", x)
})
```

```
## [[1]]
## [1] "Hello from process #38185. I see x and it is equal to 20"
##
## [[2]]
## [1] "Hello from process #38187. I see x and it is equal to 20"
##
## [[3]]
## [1] "Hello from process #38186. I see x and it is equal to 20"
##
## [[4]]
## [1] "Hello from process #38184. I see x and it is equal to 20"
```

```
# 4. STOP THE CLUSTER
stopCluster(cl)
```

## Ex 2: Parallel regressions

**Problem:** Run multiple regressions on a very wide dataset. We need to fit the following model:

$$y = X_i\beta_i + \varepsilon, \quad \varepsilon \sim N(0, \sigma_i^2), \quad \forall i$$

```
dim(X)
```

```
## [1] 500 999
```

```
X[1:6, 1:5]
```

```
##           x001           x002           x003           x004           x005
## 1  0.61827227  1.72847041 -1.4810695 -0.2471871  1.4776281
## 2  0.96777456 -0.19358426 -0.8176465  0.6356714  0.7292221
## 3 -0.04303734 -0.06692844  0.9048826 -1.9277964  2.2947675
## 4  0.84237608 -1.13685605 -1.8559158  0.4687967  0.9881953
## 5 -1.91921443  1.83865873  0.5937039 -0.1410556  0.6507415
## 6  0.59146153  0.81743419  0.3348553 -1.8771819  0.8181764
```

```
str(y)
```

```
## num [1:500] -0.8188 -0.5438 1.0209 0.0467 -0.4501 ...
```

## Ex 2: Parallel regressions (cont'd 1)

**Serial solution:** Use `apply` (forloop) to solve it

```
#  
#  
#  
ans <- apply(  
  X      = X,  
  MARGIN = 2,  
  FUN    = function(x) coef(lm(y ~ x))  
)  
ans[,1:5]
```

```
##              x001      x002      x003      x004      x005  
## (Intercept) -0.03449819 -0.03339681 -0.03728140 -0.03644192 -0.03717344  
## x           -0.06082548  0.02748265 -0.01327855 -0.08012361 -0.04067826
```

## Ex 2: Parallel regressions (cont'd 2)

**Parallel solution:** Use parApply

```
library(parallel)
cl <- makePSOCKcluster(4L)
clusterExport(cl, "y")
ans <- parApply(
  cl      = cl,
  X       = X,
  MARGIN  = 2,
  FUN     = function(x) coef(lm(y ~ x))
)
ans[,1:5]
```

```
##                x001        x002        x003        x004        x005
## (Intercept) -0.03449819 -0.03339681 -0.03728140 -0.03644192 -0.03717344
## x           -0.06082548  0.02748265 -0.01327855 -0.08012361 -0.04067826
```

Are we going any faster?

```
microbenchmark::microbenchmark(  
  parallel = parApply(  
    cl = cl,  
    X = X, MARGIN = 2,  
    FUN = function(x) coef(lm(y ~ x))  
  ),  
  serial = apply(  
    X = X, MARGIN = 2,  
    FUN = function(x) coef(lm(y ~ x))  
  ), unit = "ms"  
)
```

```
## Unit: milliseconds  
##      expr      min       lq      mean    median       uq      max  neval  
## parallel 189.3619 211.8326 226.1504 221.3169 230.7429 345.5169   100  
##      serial 553.7184 586.4854 604.3711 598.8101 609.3490 776.1760   100
```



# Extended Example: SARS-CoV2 simulation

An altered version of [Conway's game of life](#)

1. People live in torus, each individual having 8 neighbors.
2. A healthy individual interacting with a sick neighbor has the following probabilities of contracting the disease:
  - a. 100% if neither wears a face-mask.
  - b. 50% if only he wears the face-mask.
  - c. 20% if only his neighbor wears the mask.
  - d. 5% if both wear the face-mask.
3. Infected individuals may die with probability 10%.

We want to illustrate the importance of wearing face masks. We need to simulate a system with 2,500 (50 x 50) individuals, 1,000 times so we can analyze: (a) contagion curve, (b) death curve.

More models like this: The [SIRD model](#) (Susceptible-Infected-Recovered-Deceased)

# Conway's Game of Masks

Download the program [here](#).

```
source("sars-cov2.R", echo=FALSE)

# Looking at some constants
probs_sick # Sick individual's probabilities
```

```
##   deceased   infected   recovered
##      0.1         0.4         0.5
```

```
probs_susc # Probabilities of i getting the disease
```

```
##               j doesn't wear j wears
## i doesn't wear      0.9      0.20
## i wears             0.5      0.05
```

## First look: How does the simulation looks like?

```
set.seed(7123)
one <- simulate_covid(
  pop_size = 1600,
  nsick    = 160,
  nwears_mask = 1:400,
  nsteps   = 20,
  store    = TRUE
)

one$statistics[c(1:5, 16:20),]
```

##	susceptible	infected	recovered	deceased
## 0	1440	160	0	0
## 1	1265	234	85	16
## 2	1064	307	190	39
## 3	876	321	334	69
## 4	717	287	499	97
## 15	430	1	990	179
## 16	429	2	990	179
## 17	429	0	992	179
## 18	429	0	992	179
## 19	429	0	992	179


## First look: How does the simulation looks like? (contd')

```
# Location of who wears the facemask. This step is only for plotting
wears <- which(one$wears, arr.ind = TRUE) - 1
wears <- wears/(one$nr) * (1 + 1/one$nr)

# Initializing the animation
fig <- magick::image_device(600, 600, res = 96/2, pointsize = 24)
for (i in 1:one$current_step) {

  # Plot
  image(
    one$temporal[,i], col=c("gray", "tomato", "steelblue", "black"),
    main = paste("Time", i - 1L, "of", one$steps),
    xlim = c(1,4)
  )
  points(wears, col="white", pch=20, cex=1.5)
  legend(
    "topright",
    col = c("gray", "tomato", "steelblue", "black", "black"),
    legend = c(names(codes), "wears a mask"),
    pch = c(rep(15, 4), 21)
  )
}

# Finalizing plot and writing the animation
dev.off()
animation <- magick::image_animate(fig, fps = 2)
magick::image_write(animation, "covid1.gif")
```




```

set.seed(123355)
stats_nobody_wears_masks <- replicate(50, {
  simulate_covid(
    pop_size = 900,
    nsick = 10,
    nwears_mask = 0,
    nsteps = 15)$statistics[, "deceased"]
}, simplify = FALSE
)

set.seed(123355)
stats_half_wears_masks <- replicate(50, {
  simulate_covid(
    pop_size = 900,
    nsick = 10,
    nwears_mask = 450,
    nsteps = 15)$statistics[, "deceased"]
}, simplify = FALSE
)

set.seed(123355)
stats_all_wears_masks <- replicate(50, {
  simulate_covid(
    pop_size = 900,
    nsick = 10,
    nwears_mask = 900,
    nsteps = 15)$statistics[, "deceased"]
}, simplify = FALSE
)

```



Cumulative number of deceased as a function of whether none, half, or all individuals wear a face mask.

# Speed things up: Timing under the serial implementation

We will use the function `system.time` to measure how much time it takes to complete 20 simulations in serial versus parallel fashion using 4 cores

```
time_serial <- system.time({  
  ans_serial <- replicate(50, {  
    simulate_covid(  
      pop_size = 900,  
      nsick    = 10,  
      nwears_mask = 900,  
      nsteps   = 20)$statistics[, "deceased"]  
    },  
    simplify = FALSE  
  )  
})
```



# Speed things up: Parallel a Forking Cluster

Alternative 1: If you are using Unix-like system (Ubuntu, OSX, etc.), you can take advantage of process forking, and thus, parallel's `mclapply` function:

```
set.seed(1231)
time_parallel_fork <- system.time({
  ans_parallel <- parallel::mclapply(1:50, function(i) {
    simulate_covid(
      pop_size = 900,
      nsick = 10,
      nwears_mask = 900,
      nsteps = 20)$statistics[, "deceased"]
  }, mc.cores = 2L
})
```

# Speed things up: Parallel with a Socket Cluster

Alternative 2: Regardless of the operating system, we can use a Socket cluster, which is simply a group of fresh R sessions that listen to the parent/main/mother session.

```
# Step 1: Make the cluster  
cl <- parallel::makePSOCKcluster(2L)
```

```
# Step 2: Prepare the cluster  
# We could either export all the needed variables  
parallel::clusterExport(  
  cl,  
  c("calc_stats", "codes", "dat", "get_neighbors", "init", "probs_sick",  
    "probs_susc", "simulate_covid", "update_status", "update_status_all"  
)  
)
```

Or simply running the simulation script in the other sessions

```
# Step 2 (alt): Prepare the cluster  
parallel::clusterEvalQ(cl, source("sars-cov2.R"))  
parallel::clusterSetRNGStream(cl, 123) # Make sure it is reproducible!
```

```
(pids <- c(
  master = Sys.getpid(),
  offspring = unlist(parallel::clusterEvalQ(cl, Sys.getpid()))
))
#      master offspring1 offspring2
#    14810      15998      16012
```

If you are using Unix, you can see more details:

UID	PID	PPID	C	STIME	TTY	TIME	CMD
george	14810	10376	0	12:31	?	00:00:09	/usr/lib/
george	15998	1	0	12:56	?	00:00:00	/usr/lib/
george	16012	1	0	12:56	?	00:00:00	/usr/lib/

## Speed things up: Parallel with a Socket Cluster (cont'd)

```
# Step 3: Do your call
time_parallel_sock <- system.time({
  ans_parallel <- parallel::parLapply(cl, 1:50, function(i) {
    simulate_covid(
      pop_size = 900,
      nsick    = 10,
      nwears_mask = 900,
      nsteps   = 20)$statistics[, "deceased"]
    })
  })

# Step 4: Stop
parallel::stopCluster(cl)
```

Using two threads/processes, you can obtain the following speedup

```
time_serial;time_parallel_sock;time_parallel_fork
```

```
##      user  system elapsed  
## 20.514    0.238   21.166
```

```
##      user  system elapsed  
##  0.001    0.000   11.585
```

```
##      user  system elapsed  
##  0.003    0.004   11.925
```

# Cloud Computing (a.k.a. on-demand computing)

HPC clusters, super-computers, etc. need not to be bought... you can rent:

- [Amazon Web Services \(AWS\)](#)
- [Google Cloud Computing](#)
- [Microsoft Azure](#)

These services provide more than just computing (storage, data analysis, etc.). But for computing and storage, there are other free resources, e.g.:

- [The Extreme Science and Engineering Discovery Environment \(XSEDE\)](#)

# There are many ways to run R in the cloud

At USC:

- Center for Advanced Research Computing (CARC). USC users can request hundreds of cores (literally). Take a look at the [slurmR package](#)

Running R in:

- Google Cloud: <https://cloud.google.com/solutions/running-r-at-scale>
- Amazon Web Services: <https://aws.amazon.com/blogs/big-data/running-r-on-aws/>
- Microsoft Azure: <https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/r-developers-guide>

## Submitting jobs

- A key feature of cloud services > interact via command line.
- You will need to familiarize with `Rscript` and `R CMD BATCH`.
- Which is better? It depends on the application.



# Submitting jobs (examples)

Imagine we have the following R script (download [here](#)):

```
library(data.table)
set.seed(1231)
dat <- data.table(y = rnorm(1e3), x = sample.int(5, 1e3, TRUE))
dat[, mean(y), by = x]
```

## R CMD BATCH

This will run a non-interactive R session and put all the output ([stdout](#) and [stderr](#)) to the file `dummy.Rout`.

```
R CMD BATCH --vanilla dummy.R dummy.Rout &
```

## Rscript

This will also execute R in the background, with the difference that the output `dummy.Rout` will not capture `stderr` (messages, warnings and errors from R).

```
Rscript --vanilla dummy.R > dummy.Rout &
```

The `&` at the end makes sure the job is submitted and does not wait for it to end. Try it yourself (5 mins)!

# Rscript

The R script can be executed as program directly, if you specify where the `Rscript` program lives.

The following example works in Unix. This is an R script named `since_born.R` (download [here](#))

```
#!/usr/bin/Rscript
args <- tail(commandArgs(), 0)
message(Sys.Date() - as.Date(args), " days since you were born.")
```

This R script, can be executed in various ways...

## Rscript as a program

For this we would need to change it to an executable. In unix you can use the [chmod](#) command: `chmod +x since_born.R`. This allows to:

```
./since_born.R 1988-03-02
```

## Rscript in a bash script (most common)

In the case of running jobs in a cluster or something similar, we usually need to have a bash script, In our case, here we have a file named `since_born_bash.sh` that calls `Rscript` (download [here](#))

```
#!/bin/bash
Rscript since_born.R 1988-03-02
```

Which we would execute something like this

```
sh since_born_bash.sh
```

```
## 12419 days since you were born.
```

# Summary

- Parallel computing can speed up things.
- Not always needed... need to make sure that you are taking advantage of vectorization.
- Most of the time we look at "Embarrassingly parallel computing."
- In R, explicit parallelism can be achieved using the **parallel** package:
  1. Load the package and create a cluster **parallel::makeCluster()**
  2. Setup the environment **parallel::clusterEvalQ()**, **parallel::clusterExport()**, and **parallel::clusterSetRNGStream()**
  3. Make the call, e.g., **parallel::parLapply()**
  4. Stop the cluster **parallel::stopCluster()**
- Regardless of the Cloud computing service we are using, we will be using either R CMD BATCH or Rscript to submit jobs.

# Session info

```
## R version 4.1.2 (2021-11-01)
## Platform: x86_64-apple-darwin17.0 (64-bit)
## Running under: macOS Big Sur 10.16
##
## Matrix products: default
## BLAS:   /Library/Frameworks/R.framework/Versions/4.1/Resources/lib/libRblas.0.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/4.1/Resources/lib/libRlapack.dylib
##
## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## attached base packages:
## [1] parallel stats      graphics  grDevices utils      datasets  methods
## [8] base
##
## loaded via a namespace (and not attached):
## [1] digest_0.6.29  R6_2.5.1      jsonlite_1.7.2 magrittr_2.0.2
## [5] evaluate_0.14  highr_0.9      xaringan_0.22  stringi_1.7.6
## [9] rlang_1.0.0     cli_3.1.1      rstudioapi_0.13 jquerylib_0.1.4
## [13] bslib_0.3.1     rmarkdown_2.11 tools_4.1.2     stringr_1.4.0
## [17] xfun_0.29       yaml_2.2.1     fastmap_1.1.0  compiler_4.1.2
## [21] htmltools_0.5.2 knitr_1.37     sass_0.4.0
```


# Resources

- [Package parallel](#)
- [Using the iterators package](#)
- [Using the foreach package](#)
- [32 OpenMP traps for C++ developers](#)
- [The OpenMP API specification for parallel programming](#)
- ['openmp' tag in Rcpp gallery](#)
- [OpenMP tutorials and articles](#)

For more, checkout the [CRAN Task View on HPC](#){target="\_blank"}

# Simulating $\pi$

- We know that  $\pi = \frac{A}{r^2}$ . We approximate it by randomly adding points  $x$  to a square of size 2 centered at the origin.
- So, we approximate  $\pi$  as  $\Pr\{\|x\| \leq 1\} \times 2^2$





The R code to do this

```
pisim <- function(i, nsim) { # Notice we don't use the -i-  
  # Random points  
  ans <- matrix(runif(nsim*2), ncol=2)  
  
  # Distance to the origin  
  ans <- sqrt(rowSums(ans^2))  
  
  # Estimated pi  
  (sum(ans <= 1)*4)/nsim  
}
```

```

library(parallel)
# Setup
cl <- makePSOCKcluster(4L)
clusterSetRNGStream(cl, 123)

# Number of simulations we want each time to run
nsim <- 1e5

# We need to make -nsim- and -pisim- available to the
# cluster
clusterExport(cl, c("nsim", "pisim"))

# Benchmarking: parSapply and sapply will run this simulation
# a hundred times each, so at the end we have 1e5*100 points
# to approximate pi
microbenchmark::microbenchmark(
  parallel = parSapply(cl, 1:100, pisim, nsim=nsim),
  serial   = sapply(1:100, pisim, nsim=nsim), times = 1, unit="ms"
)

```

## Unit: milliseconds

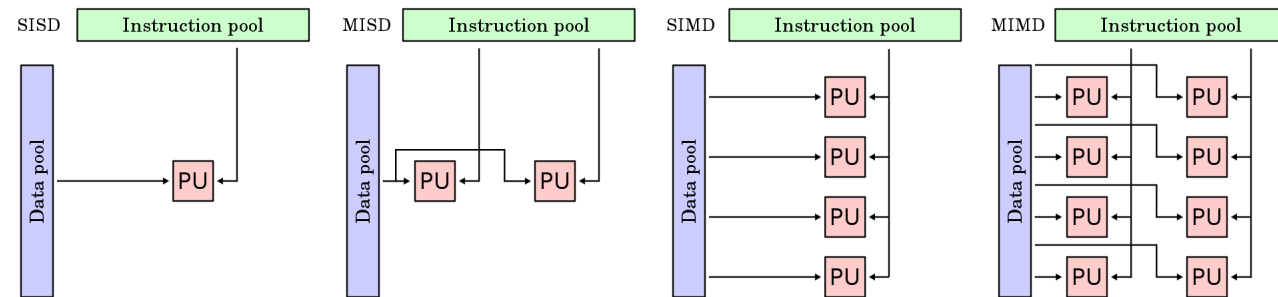
##	expr	min	lq	mean	median	uq	max	neval
##	parallel	320.6614	320.6614	320.6614	320.6614	320.6614	320.6614	1
##	serial	861.9523	861.9523	861.9523	861.9523	861.9523	861.9523	1

## (Bonus) Overview of HPC

Using [Flynn's classical taxonomy](#), we can classify parallel computing according to the following two dimensions:

a. Type of instruction: Single vs Multiple

b. Data stream: Single vs Multiple



[Michael Flynn's](#) Taxonomy ([wiki](#))

## (Bonus) Parallel computing: Software

Implicit parallelization:

- [tensorflow](#): Machine learning framework
- [pqR](#): Branched version of R.
- [Microsoft R](#): Microsoft's R private version (based on Revolution Analytics' R version).
- [data.table](#) (R package): Data wrangling using multiple cores.
- [caret](#) (R package): A meta package, has various implementations using parallel computing.

Explicit parallelization ([DIY](#)):

- [CUDA](#) (C/C++ library): Programming with GP-GPUs.
- [Open MP](#) (C/C++ library): Multi-core programming (CPUs).
- [Open MPI](#) (C/C++ library): Large scale programming with multi-node systems.
- [Threading Building Blocks](#) (C/C++ library): Intel's parallel computing library.
- [Kokkos](#) (C++ library): A hardware-agnostic programming framework for HPC applications.
- [parallel](#) (R package): R's built-in parallel computing package
- [future](#) (R package): Framework for parallelizing R.
- [RcppParallel](#) (R C++ API wrapper): Header and templates for building [Rcpp](#)+multi-threaded programs.
- [julia](#) (programming language): High-performing, has a framework for parallel computing as well.